

Project: Data Lake with EMR (Spark & Co. @ AWS)

Creating EMR Cluster ("emr-5.31.0" - Spark 2.4.6 , Zeppelin 0.8.2 , Hadoop 2.10.0) of AWS

This last EMR Version does have the correct setup up with JupyterEnterpriseGateway !

Overview and Steps

- Step 0: Create Base Config (Setup EMR Cluster (with Spark), Create Notebook, Attach Notebook to Cluster)
- Step 1: Create Spark Session
- Step 2: Define Data Paths and Access Data
- Step 3: Process Song Data (read out, define schemas, write as parquet)
- Step 4: Process Log Data (read out, define schemas, write as parquet)
- Step 5: Run [et1.py] script with clean code
- Step 6: Clean up Resources (EMR, Notebook, S3)

Version 3 (AWS, EMR) - Revision 09 - 2023/04/04 Mr Morphy - [GitHub Profile \(https://github.com/MrMorphy\)](https://github.com/MrMorphy)

GitHub Project - [udacity-course-proj-data-lake \(https://github.com/MrMorphy/udacity-course-proj-data-lake\)](https://github.com/MrMorphy/udacity-course-proj-data-lake)

Step 0: Create Base Config

- Setup EMR Cluster (with Spark) at AWS UI
- Create Notebook at AWS UI
- Attach Notebook to 'Waiting' EMR Cluster

Load libraries to execute code

```
In [1]: #import configparser
from datetime import datetime
import os
```

Starting Spark application

ID	YARN Application ID	Kind	State	Spark UI
0	application_1680641614700_0001	pyspark	idle	Link (http://ip-172-31-5-17.us-west-2.compute.internal:20888/proxy/application_1680641614700_0001/)

SparkSession available as 'spark'.

```
In [2]: from pyspark.sql import SparkSession
```

```
In [3]: from pyspark.sql.functions import udf, col, monotonically_increasing_id
from pyspark.sql.functions import year, month, dayofmonth, hour, weekofyear, date_format
```

```
In [4]: from pyspark.sql.types import StructType as R, StructField as Fld, \
        DoubleType as Db1, StringType as Str, \
        IntegerType as Int, DateType as Date, TimestampType
```

Step 1: Create Spark Session

```
In [6]: # DEBUG
spark
```

<pyspark.sql.session.SparkSession object at 0x7fc0cb14bdd0>

Step 2: Define Data Paths and Access Data

```
In [7]: # AWS: Online @ AWS - S3:
input_data = "s3a://udacity-dend/"
output_data = "s3a://data-lake-project-out/"
```

Step 3: Process Song Data

(read out, define schemas, write as parquet)

```
In [8]: # AWS:
#def process_song_data(spark, input_data, output_data):

# get filepath to FULL SET of song data file [song_data/A/B/C/TRABCEI128F424C983.json]
song_data = input_data + 'song_data/**/*.json'
```

```
In [9]: # define the song schema (like [staging_songs] table)
songSchema = R([
    Fld("artist_id",      Str()),
    Fld("artist_latitude", Dbl()),
    Fld("artist_location", Str()),
    Fld("artist_longitude", Dbl()),
    Fld("artist_name",     Str()),
    Fld("duration",        Dbl()),
    Fld("num_songs",        Int()),
    Fld("song_id",         Str()),
    Fld("title",           Str()),
    Fld("year",            Int()),
])
```

```
In [10]: # read song data JSON file into data frame
df = spark.read.json(song_data, schema=songSchema)
```

► Spark Job Progress

```
In [11]: # DEBUG
# songSchema given
df.printSchema()
```

```
root
|-- artist_id: string (nullable = true)
|-- artist_latitude: double (nullable = true)
|-- artist_location: string (nullable = true)
|-- artist_longitude: double (nullable = true)
|-- artist_name: string (nullable = true)
|-- duration: double (nullable = true)
|-- num_songs: integer (nullable = true)
|-- song_id: string (nullable = true)
|-- title: string (nullable = true)
|-- year: integer (nullable = true)
```

```
In [13]: # df.count()
print('>> [' + str(df.count()) + '] songs from song_data read out in JSON-format')
```

► Spark Job Progress

```
>> [14896] songs from song_data read out in JSON-format
```

```
In [12]: # DEBUG
df.show(5)
```

► Spark Job Progress

```
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+
|      artist_id|artist_latitude|      artist_location|artist_longitude|      artist_name| dur
ation|num_songs|      song_id|      title|year|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+
|AR4T2IF1187B9ADB87|      63.96027|<a href="http://b...|      10.22442|      Billy Idol|233.
22077|      1|SOVIYJY12AF72A4B00|The Dead Next Doo...|1983|
|AR4T2IF1187B9ADB87|      63.96027|<a href="http://b...|      10.22442|      Billy Idol|287.
92118|      1|SOVYXYL12AF72A3373|Rebel Yell (1999 ...|1983|
|ARQ846I1187B9A7083|      null|      |      null|Yvonne S. Moriart...|196.
04853|      1|SOEPTVC12A67ADD0DA|To Zucchabar ["Gl...| 0|
|AR4T2IF1187B9ADB87|      63.96027|<a href="http://b...|      10.22442|      Billy Idol|247.
53587|      1|SOLQYSZ12AB0181F97|      Mony Mony (Live)|1987|
|AR3TZ691187FB3DBB1|      null|      |      null|Russell Watson / ...|273.
44934|      1|SOVPFJK12A6701CB16|Barcelona - (Frie...|2000|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+
only showing top 5 rows
```

```
In [14]: # DEBUG
df1 = df.filter(df.title == 'Young Boy Blues')
df1.show(5)
```

► Spark Job Progress

```
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+
|artist_id|artist_latitude|artist_location|artist_longitude|artist_name|duration|num_songs|song_id|
title|year|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+
```

Working with only a subset of the files, to save time!

```
In [23]: # get filepath to song data file [song_data/A/B/C/TRABCEI128F424C983.json]
#song_data = input_data + 'song_data/**/*.json'
# Experiment with a subset of data:
song_data = input_data + 'song_data/A/A/A/*.json'
```

```
In [24]: # read song data JSON file into data frame
df = spark.read.json(song_data, schema=songSchema)
```

```
In [25]: # df.count()
print('>> [' + str(df.count()) + '] songs from song_data read out in JSON-format')
```

► Spark Job Progress

```
>> [24] songs from song_data read out in JSON-format
```

```
In [26]: # extract columns to create songs table
song_columns = ["song_id", "title", "artist_id", "year", "duration"]
songs_table = df.select(song_columns) \
    .dropDuplicates()
```

```
In [16]: # DEBUG
songs_table.printSchema()
```

```
root
|-- song_id: string (nullable = true)
|-- title: string (nullable = true)
|-- artist_id: string (nullable = true)
|-- year: integer (nullable = true)
|-- duration: double (nullable = true)
```

```
In [17]: # DEBUG
songs_table.show(5)
```

► Spark Job Progress

song_id	title	artist_id	year	duration
SOIIRYF12AAF3B4858	Pitbulls in the P...	ARZNUT51187FB44098	0	230.29506
SOGUKZL12A6D4FDF28	The Apples Are Ju...	ARY1NR31187FB45277	0	156.39465
SORXLTC12A81C20B18	La Chiave (feat.S...	ARLN9UN1187FB481FD	0	198.00771
SOAIWKG12AB018325D	Bach Prelude In G...	ARTPL6J1187FB4D6FB	1992	136.33261
SOFZFSD12A8C13DEF4	Sweet Sophia (Alb...	ARYPW621187FB5B49E	2007	191.76444

only showing top 5 rows

```
In [28]: # [X] VERIFICAR @ AWS...
# write songs table to parquet files partitioned by year and artist
songs_table.repartition('year', 'artist_id') \
    .write.partitionBy('year', 'artist_id') \
    .parquet(output_data + 'songs_data/songs_table.parquet', 'overwrite')
```

► Spark Job Progress

```
In [ ]:
```

```
In [29]: # extract columns to create artists table
artists_column = ["artist_id", "artist_name as name", "artist_location as location", \
    "artist_latitude as latitude", "artist_longitude as longitude"]
artists_table = df.selectExpr(artists_column).dropDuplicates()
```

```
In [30]: # DEBUG
artists_table.printSchema()
```

```
root
|-- artist_id: string (nullable = true)
|-- name: string (nullable = true)
|-- location: string (nullable = true)
|-- latitude: double (nullable = true)
|-- longitude: double (nullable = true)
```

```
In [31]: # DEBUG
artists_table.show(5)
```

► Spark Job Progress

artist_id	name	location	latitude	longitude
ARZKCQM1257509D107	Dataphiles		null	null
ARKYKXP11F50C47A6A	The Supersuckers		null	null
ARGE7G11187FB37E05	Cyndi Lauper	Brooklyn, NY	null	null
ARY589G1187B9A9F4E	Talkdemonic	Portland, OR	45.51179	-122.67563
ARA23XO1187B9AF18F	The Smithereens	Carteret, New Jersey	40.57885	-74.21956

only showing top 5 rows

```
In [32]: # write artists table to parquet files
artists_table.write.parquet(output_data + 'artists_data/artists_table.parquet', 'overwrite')
```

► Spark Job Progress

Step 4: Process Log Data

(read out, define schemas, write as parquet)

```
In [33]: # AWS:
#def process_log_data(spark, input_data, output_data):

# get filepath to FULL Data Set Log data file [log_data/2018/11/2018-11-17-events.json]
log_data = input_data + 'log_data/**/*.json'
```

```
In [34]: # read Log data file into data frame
df = spark.read.json(log_data)
```

► Spark Job Progress

```
In [35]: # number of Lines
# df.count()
print('>> [' + str(df.count()) + '] logs entries read IN, of JSON logs_data')
```

► Spark Job Progress

>> [8056] logs entries read IN, of JSON logs_data

```
In [36]: # DEBUG
df.printSchema()
```

```
root
|-- artist: string (nullable = true)
|-- auth: string (nullable = true)
|-- firstName: string (nullable = true)
|-- gender: string (nullable = true)
|-- itemInSession: long (nullable = true)
|-- lastName: string (nullable = true)
|-- length: double (nullable = true)
|-- level: string (nullable = true)
|-- location: string (nullable = true)
|-- method: string (nullable = true)
|-- page: string (nullable = true)
|-- registration: double (nullable = true)
|-- sessionId: long (nullable = true)
|-- song: string (nullable = true)
|-- status: long (nullable = true)
|-- ts: long (nullable = true)
|-- userAgent: string (nullable = true)
|-- userId: string (nullable = true)
```

Working with only a subset of the files, to save time!

```
In [37]: # 1) experiment with a subset of the files,
# 2) and with local data (extract before from /data/log-data.zip)
log_data = input_data + 'log_data/2018/11/*.json'
```

```
In [38]: # read Log data file into data frame
df = spark.read.json(log_data)
```

► Spark Job Progress

```
In [39]: # number of Lines
# df.count()
print('>> [' + str(df.count()) + '] logs entries read IN, of JSON logs_data')
```

► Spark Job Progress

>> [8056] logs entries read IN, of JSON logs_data

```
In [40]: # DEBUG
df.show(5)
```

► Spark Job Progress

```
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+
|   artist|   auth|firstName|gender|itemInSession|lastName|  length|level|          location
|method|   page|   registration|sessionId|          song|status|          ts|          userAg
ent|userId|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+
|  Harmonia|Logged In|   Ryan|  M|          0|  Smith|655.77751| free|San Jose-Sunnyval...
| PUT|NextSong|1.541016707796E12|          583|  Sehr kosmisch|  200|1542241826796|"Mozilla/5.0 (X1
1...|  26|
|The Prodigy|Logged In|   Ryan|  M|          1|  Smith|260.07465| free|San Jose-Sunnyval...
| PUT|NextSong|1.541016707796E12|          583|The Big Gundown|  200|1542242481796|"Mozilla/5.0 (X1
1...|  26|
|   Train|Logged In|   Ryan|  M|          2|  Smith|205.45261| free|San Jose-Sunnyval...
| PUT|NextSong|1.541016707796E12|          583|   Marry Me|  200|1542242741796|"Mozilla/5.0 (X1
1...|  26|
|      null|Logged In|  Wyatt|  M|          0|  Scott|      null| free|Eureka-Arcata-For...
| GET|   Home|1.540872073796E12|          563|      null|  200|1542247071796|Mozilla/5.0 (Win
d...|  9|
|      null|Logged In|  Austin|  M|          0| Rosales|      null| free|New York-Newark-J...
| GET|   Home|1.541059521796E12|          521|      null|  200|1542252577796|Mozilla/5.0 (Win
d...| 12|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+
only showing top 5 rows
```

```
In [41]: # filter by actions for song plays
df = df.filter(df.page == 'NextSong')
```

```
In [42]: # number of reduce amount of lines
# df.count()
print('>> [' + str(df.count()) + '] songs filtered "NextSong" of logs_data')
```

► Spark Job Progress

>> [6820] songs filtered "NextSong" of logs_data

```
In [43]: # extract columns for USERS table
# artists_table = # TYPO at resources (*.zip)!
users_columns = ["userId as user_id", "firstName as first_name", \
                 "lastName as last_name", "gender", "level"]
users_table = df.selectExpr(users_columns).dropDuplicates()
```

```
In [44]: # DEBUG
users_table.printSchema()
```

```
root
 |-- user_id: string (nullable = true)
 |-- first_name: string (nullable = true)
 |-- last_name: string (nullable = true)
 |-- gender: string (nullable = true)
 |-- level: string (nullable = true)
```

```
In [45]: # DEBUG
users_table.show(5)
```

► Spark Job Progress

```
+-----+-----+-----+-----+-----+
|user_id|first_name|last_name|gender|level|
+-----+-----+-----+-----+-----+
|      81|      Sienna|      Colon|      F| free|
|      75|      Joseph|Gutierrez|      M| free|
|      47|      Kimber|      Norris|      F| free|
|      20|      Aiden|      Ramirez|      M| paid|
|      90|      Andrea|      Butler|      F| free|
+-----+-----+-----+-----+-----+
only showing top 5 rows
```

```
In [46]: # write USERS table to parquet files
users_table.write.parquet(output_data + 'users_data/users_table.parquet', 'overwrite')
```

► Spark Job Progress

```
In [47]: # create "timestamp" column from original timestamp ("ts") column
get_datetime = udf(lambda x: str(datetime.fromtimestamp(int(x)/1000.0)))
dfTimestamp = df.withColumn("start_time", get_datetime(df.ts))
```

```
In [48]: # extract columns to create time table
time_table = dfTimestamp.select("start_time").dropDuplicates() \
    .withColumn("hour", hour(col("start_time"))) \
    .withColumn("day", dayofmonth(col("start_time"))) \
    .withColumn("week", weekofyear(col("start_time"))) \
    .withColumn("month", month(col("start_time"))) \
    .withColumn("year", year(col("start_time"))) \
    .withColumn("weekday", date_format(col("start_time"), 'E'))
```

```
In [50]: # DEBUG
time_table.printSchema()
time_table.show(5)
```

► Spark Job Progress

```
root
|-- start_time: string (nullable = true)
|-- hour: integer (nullable = true)
|-- day: integer (nullable = true)
|-- week: integer (nullable = true)
|-- month: integer (nullable = true)
|-- year: integer (nullable = true)
|-- weekday: string (nullable = true)

+-----+-----+-----+-----+-----+-----+
| start_time | hour | day | week | month | year | weekday |
+-----+-----+-----+-----+-----+-----+
| 2018-11-15 07:51:... | 7 | 15 | 46 | 11 | 2018 | Thu |
| 2018-11-15 10:53:... | 10 | 15 | 46 | 11 | 2018 | Thu |
| 2018-11-21 09:13:... | 9 | 21 | 47 | 11 | 2018 | Wed |
| 2018-11-15 17:10:... | 17 | 15 | 46 | 11 | 2018 | Thu |
| 2018-11-15 14:34:... | 14 | 15 | 46 | 11 | 2018 | Thu |
+-----+-----+-----+-----+-----+-----+

only showing top 5 rows
```

```
In [49]: # write time table to parquet files partitioned by year and month
time_table.repartition('year', 'month') \
    .write.partitionBy("year", "month") \
    .parquet(output_data + 'time_data/time_table.parquet', 'overwrite')
```

► Spark Job Progress

Preparation and Code for songplays_table

```
In [52]: get_datetime = udf(lambda x: str(datetime.fromtimestamp( int(x)/1000.0 )) )
df = df.withColumn("start_time", get_datetime(df.ts))
df = df.withColumn("year", year(col("start_time")))
df = df.withColumn("month", month(col("start_time")))
```



```
In [55]: # DEBUG
df.printSchema()
df.show(5)
```

► Spark Job Progress

```
root
|-- artist: string (nullable = true)
|-- auth: string (nullable = true)
|-- firstName: string (nullable = true)
|-- gender: string (nullable = true)
|-- itemInSession: long (nullable = true)
|-- lastName: string (nullable = true)
|-- length: double (nullable = true)
|-- level: string (nullable = true)
|-- location: string (nullable = true)
|-- method: string (nullable = true)
|-- page: string (nullable = true)
|-- registration: double (nullable = true)
|-- sessionId: long (nullable = true)
|-- song: string (nullable = true)
|-- status: long (nullable = true)
|-- ts: long (nullable = true)
|-- userAgent: string (nullable = true)
|-- userId: string (nullable = true)
|-- start_time: string (nullable = true)
|-- year: integer (nullable = true)
|-- month: integer (nullable = true)

+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|   artist|   auth|firstName|gender|itemInSession|lastName|  length|level|      location
|method|   page|  registration|sessionId|              song|status|      ts|      u
serAgent|userId|      start_time|year|month|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|  Harmonia|Logged In|   Ryan|   M|          0|  Smith|655.77751| free|San Jose-Sunnyval...
|  PUT|NextSong|1.541016707796E12|    583|      Sehr kosmisch|   200|1542241826796|"Mozilla/5.0
(X11...|   26|2018-11-15 00:30:...|2018|   11|
|The Prodigy|Logged In|   Ryan|   M|          1|  Smith|260.07465| free|San Jose-Sunnyval...
|  PUT|NextSong|1.541016707796E12|    583|  The Big Gundown|   200|1542242481796|"Mozilla/5.0
(X11...|   26|2018-11-15 00:41:...|2018|   11|
|   Train|Logged In|   Ryan|   M|          2|  Smith|205.45261| free|San Jose-Sunnyval...
|  PUT|NextSong|1.541016707796E12|    583|    Marry Me|   200|1542242741796|"Mozilla/5.0
(X11...|   26|2018-11-15 00:45:...|2018|   11|
|Sony Wonder|Logged In| Samuel|   M|          0|Gonzalez|218.06975| free|Houston-The Woodl...
|  PUT|NextSong|1.540492941796E12|    597|    Blackbird|   200|1542253449796|"Mozilla/5.0
(Mac...|   61|2018-11-15 03:44:...|2018|   11|
|  Van Halen|Logged In|  Tegan|   F|          2|  Levine|289.38404| paid|Portland-South Po...
|  PUT|NextSong|1.540794356796E12|    602|Best Of Both Worl...|   200|1542260935796|"Mozilla/5.0
(Mac...|   80|2018-11-15 05:48:...|2018|   11|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
only showing top 5 rows
```

```
In [56]: # read in song data to use for songplays table
song_df = spark.read.option("mergeSchema", "true") \
               .parquet(output_data + "songs_data/songs_table.parquet")
```

► Spark Job Progress

```
In [57]: # DEBUG
print('>> [' + str(song_df.count()) + '] songs readout from songs_table.PARQUET')
```

► Spark Job Progress

```
>> [24] songs readout from songs_table.PARQUET
```

```
In [58]: # DEBUG
song_df.printSchema()
song_df.show(5)
```

► Spark Job Progress

```
root
|-- song_id: string (nullable = true)
|-- title: string (nullable = true)
|-- duration: double (nullable = true)
|-- year: integer (nullable = true)
|-- artist_id: string (nullable = true)

+-----+-----+-----+-----+-----+
| song_id| title| duration| year| artist_id|
+-----+-----+-----+-----+-----+
|SOKTJDS12AF72A25E5|Drown In My Own T...| 192.522| 0|ARA23X01187B9AF18F|
|SOEKAZG12AB018837E|I'll Slap Your Fa...|129.85424|2001|ARSVTNL1187B992A91|
|SOAFBCP12A8C13CC7D|King Of Scurf (20...|301.40036|1972|ARTC1LV1187B9A4858|
|SORRNOC12AB017F52B|The Last Beat Of ...|337.81506|2004|ARSZ7L31187FB4E610|
|SOQPWCR12A6D4FB2A3|A Poor Recipe For...|118.07302|2005|AR73AI01187B9AD57B|
+-----+-----+-----+-----+-----+
only showing top 5 rows
```

```
In [64]: # DEBUG - JOIN - Testing (1/3)
df1 = song_df.filter(song_df.title == 'Scream')
df1.show(5)
```

► Spark Job Progress

```
+-----+-----+-----+-----+-----+
| song_id| title| duration| year| artist_id|
+-----+-----+-----+-----+-----+
|SOBLFFE12AF72AA5BA|Scream|213.9424|2009|ARJNIUY12298900C91|
+-----+-----+-----+-----+-----+
```

```
In [65]: # DEBUG - JOIN - Testing (2/3)
df2 = df.filter(df.song == 'Scream') # Riverside / Young Boy Blues / Setanta matins
df2.show(5)
```

► Spark Job Progress

```
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| artist| auth| firstName| gender| itemInSession| lastName| length| level| locati
on| method| page| registration| sessionId| song| status| ts| userAgent| use
rId| start_time| year| month|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|Michael Jackson|Logged In| Kate| F| 32| Harrell|278.282| paid|Lansing-East Lan
s...| PUT|NextSong|1.540472624796E12| 605|Scream| 200|1542298745796|"Mozilla/5.0 (X11...|
97|2018-11-15 16:19:...|2018| 11|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
```

```
In [63]: # DEBUG - JOIN - Testing (3/3)
df3 = df.join(song_df, song_df.title == df.song ).select(df.song.alias("song (df)"), song_df.title.a
df3.count()
df3.show(5)
```

► Spark Job Progress

```
+-----+-----+
|song (df)|title (song_df)|
+-----+-----+
|   Scream|          Scream|
+-----+-----+
```

```
In [66]: # extract columns from joined song and Log datasets to create songplays table
songplays_table = song_df.join(df, (song_df.title == df.song) ) \
    .select('start_time', \
            df.year, \
            df.month, \
            col('userId').alias("user_id"), \
            df.level, \
            song_df.song_id, \
            song_df.artist_id, \
            col('sessionId').alias("session_id"), \
            df.location, \
            col('userAgent').alias("user_agent") \
            )
```

```
In [67]: songplays_table = songplays_table.withColumn("songplay_id", monotonically_increasing_id())
```

```
In [68]: # DEBUG
print('>> [' + str(songplays_table.count()) + '] songs found, on JOIN matching for songsplays_table')
```

► Spark Job Progress

```
>> [1] songs found, on JOIN matching for songsplays_table
```

```
In [69]: # DEBUG
songplays_table.printSchema()
songplays_table.show(3)
```

► Spark Job Progress

An error was encountered:
Error sending http request and maximum retry encountered.

```
In [ ]: # write songplays table to parquet files partitioned by year and month
songplays_table.repartition('year', 'month') \
    .write.partitionBy("year", "month") \
    .parquet(output_data + 'songplays_data/songplays_table.parquet', 'overwrite')
```

```
In [ ]: print('>> END! (main)')
```

Step 5: Run [etl.py] script

... with working clean code

```
In [ ]: !python 'etl.py' # to be executed at the Jupyter Notebook on AWS!
```

Step 6: Clean up Resources

(EMR, Notbook, S3)

- Export Jupyter Notebook
- Terminate EMR Cluster

- Delete S3 Bucket

===== [EOF] =====