

Project: Data Lake with EMR (Spark & Co. @ AWS)

Creating EMR Cluster ("emr-5.31.0" - Spark 2.4.6 , Zeppelin 0.8.2 , Hadoop 2.10.0) of AWS
This last EMR Version does have the correct setup up with JupyterEnterpriseGateway !

Overview and Steps

- Step 0: Create Base Config (Setup IAM-Role [dl.cfg], Setup EMR Cluster (with Spark), Create Notebook)
- Step 1: Create Spark Session
- Step 2: Define Data Paths and Access Data
- Step 3: Process Song Data (read out, define schemas, write as parquet)
- Step 4: Process Log Data (read out, define schemas, write as parquet)
- Step 5: Run [etl.py] script with clean code
- Step 6: Clean up Resources (EMR, Notebook, IAM-Role)

Clear fields at [dl.cfg] before submitting project for review!

KEY = `YOUR_AWS_KEY` SECRET = `YOUR_AWS_SECRET`

Version 2 (LOC) - Revision 06 - 2023/03/20 Mr Morphy - [GitHub Profile \(https://github.com/MrMorphy\)](https://github.com/MrMorphy)

GitHub Project - [udacity-course-proj-data-lake \(https://github.com/MrMorphy/udacity-course-proj-data-lake\)](https://github.com/MrMorphy/udacity-course-proj-data-lake)

Step 0: Create Base Config

- Setup IAM-Role [dl.cfg]
- Setup EMR Cluster (with Spark) at AWS UI
- Create Notebook at AWS UI

Load libraries to execute code

In [1]:

```
import configparser
from datetime import datetime
import os
```

In [2]:

```
from pyspark.sql import SparkSession
```

In [3]:

```
from pyspark.sql.functions import udf, col, monotonically_increasing_id
from pyspark.sql.functions import year, month, dayofmonth, hour, weekofyear, date_format
```

In [4]:

```
from pyspark.sql.types import StructType as R, StructField as Fld, \
    DoubleType as Dbl, StringType as Str, \
    IntegerType as Int, DateType as Date, TimestampType
```

In [5]:

```
# only for jupyter workbook tests necessary
import pandas as pd
```

In [21]:

```
# AWS: behaviour? Pending!
config = configparser.ConfigParser()
config.read('dl.cfg')
print('>> Read Out Config-Infos from [dl.cfg]')

os.environ['AWS_ACCESS_KEY_ID'] = config['AWS_ACCESS_KEY_ID']
os.environ['AWS_SECRET_ACCESS_KEY'] = config['AWS_SECRET_ACCESS_KEY']
```

```
-----
MissingSectionHeaderError                                Traceback (most recent call last)
<ipython-input-21-73b7fd6b2e87> in <module>()
      1 config = configparser.ConfigParser()
----> 2 config.read('dl.cfg')
      3 print('>> Read Out Config-Infos from [dl.cfg]')
      4
      5 os.environ['AWS_ACCESS_KEY_ID'] = config['AWS_ACCESS_KEY_ID']

/opt/conda/lib/python3.6/configparser.py in read(self, filenames, encoding)
    695         try:
    696             with open(filename, encoding=encoding) as fp:
--> 697                 self._read(fp, filename)
    698         except OSError:
    699             continue

/opt/conda/lib/python3.6/configparser.py in _read(self, fp, fpname)
    1078         # no section header in the file?
    1079         elif cursect is None:
-> 1080             raise MissingSectionHeaderError(fpname, lineno,
line)
    1081         # an option line?
    1082         else:
```

```
MissingSectionHeaderError: File contains no section headers.
file: 'dl.cfg', line: 1
'AWS_ACCESS_KEY_ID=AKIAQ4TUNEC3HLIPRATG\n'
```

INFO: Add [AWS] on dl.cfg File as well at code " =config['AWS'][...]"

In [6]:

```
config = configparser.ConfigParser()
config.read('dl.cfg')
print('>> Read Out Config-Infos from [dl.cfg]')

os.environ['AWS_ACCESS_KEY_ID'] = config['AWS']['AWS_ACCESS_KEY_ID']
os.environ['AWS_SECRET_ACCESS_KEY'] = config['AWS']['AWS_SECRET_ACCESS_KEY']

>> Read Out Config-Infos from [dl.cfg]
```

In []:

```
# AWS: only for Jupyter Notebook execution @ AWS, but NOT @ [etl.py] required
os.environ['AWS_ACCESS_KEY_ID'] = '<AWS_ACCESS_KEY_ID>'
os.environ['AWS_SECRET_ACCESS_KEY'] = '<AWS_SECRET_ACCESS_KEY>'
```

Step 1: Create Spark Session

In [7]:

```
print('>> START (main)')
```

>> START (main)

In [8]:

```
#def create_spark_session():
spark = SparkSession \
    .builder \
    .config("spark.jars.packages", "org.apache.hadoop:hadoop-aws:2.10.0") \
    .getOrCreate()
#return spark
# @ EMR (AWS): "emr.5.20.0" > "hadoop 2.7.0"
# @ EMR (AWS): "emr-5.31.0" > "hadoop 2.10.0"
print('>> spark session created')
```

>> spark session created

In [9]:

```
# DEBUG
spark
```

Out[9]:

SparkSession - in-memory**SparkContext**Spark UI (<http://202c6ecc0754:4041>)**Version**

v2.4.3

Master

local[*]

AppName

pyspark-shell

Step 2: Define Data Paths and Access Data

In [10]:

```
# AWS: Online @ AWS - S3:
input_data = "s3a://udacity-dend/"
output_data = "s3a://udacity-project-data-lake-sparkify/output/"
```

In [10]:

```
# Local data ./data/A/A/A/*.json
input_data = "data/"
output_data = "data_output/"
```

Step 3: Process Song Data

(read out, define schemas, write as parquet)

In [11]:

```
print('>> processing song data')
```

```
>> processing song data
```

In [12]:

```
# AWS:
#def process_song_data(spark, input_data, output_data):

# get filepath to song data file [song_data/A/B/C/TRABCEI128F424C983.json]
song_data = input_data + 'song_data/**/*.json'
```

```
-----
NameError                                Traceback (most recent call last)
<ipython-input-12-4a56d3830b03> in <module>()
      2
      3 # get filepath to song data file [song_data/A/B/C/TRABCEI128F424C98
3.json]
----> 4 song_data = input_data + 'song_data/**/*.json'
```

NameError: name 'input_data' is not defined

In [23]:

```
# 1) experiment with a subset of the files,
# 2) and with local data (extract before from /data/song-data.zip)
song_data_loc = input_data + 'song_data/A/A/A/*.json'
```

In [28]:

```
# EINMALIG!
# 3) for df.join(df.song == song_df.title) ALL required
# song_data_loc = input_data_local + 'song-data.zip'
import zipfile
with zipfile.ZipFile("data/song-data.zip") as zip_ref:
    zip_ref.extractall("data/song_data_unzipped/")
# /data/song_data_unzipped/song_data/A/[A|B]/[A|B|C]/*.json
```

In [12]:

```
# get filepath to song data file [song_data/A/B/C/TRABCEI128F424C983.json]
song_data = input_data + 'song_data_unzipped/song_data/**/*.json'
```

In [13]:

```
# define the song schema (like [staging_songs] table)
songSchema = R([
    Fld("artist_id",      Str()),
    Fld("artist_latitude", Db1()),
    Fld("artist_location", Str()),
    Fld("artist_longitude", Db1()),
    Fld("artist_name",    Str()),
    Fld("duration",       Db1()),
    Fld("num_songs",      Int()),
    Fld("song_id",        Str()),
    Fld("title",          Str()),
    Fld("year",           Int()),
])
```

In [18]:

```
# AWS:  
# read song data JSON file into data frame  
df = spark.read.json(song_data, schema=songSchema)  
# ERR, maybe because of S3FullAccess.
```

```

-----
Py4JJavaError                                Traceback (most recent call last)
<ipython-input-18-68542580ee9a> in <module>()
      1 # read song data JSON file into data frame
----> 2 df = spark.read.json(song_data, schema=songSchema)

/opt/spark-2.4.3-bin-hadoop2.7/python/pyspark/sql/readwriter.py in json(self, path, schema, primitivesAsString, prefersDecimal, allowComments, allowUnquotedFieldNames, allowSingleQuotes, allowNumericLeadingZero, allowBackslashEscapingAnyCharacter, mode, columnNameOfCorruptRecord, dateFormat, timestampFormat, multiline, allowUnquotedControlChars, lineSep, samplingRatio, dropFieldIfAllNull, encoding)
    272         path = [path]
    273         if type(path) == list:
--> 274             return self._df(self._jreader.json(self._spark._sc._jvm.PythonUtils.toSeq(path)))
    275         elif isinstance(path, RDD):
    276             def func(iterator):

/opt/spark-2.4.3-bin-hadoop2.7/python/lib/py4j-0.10.7-src.zip/py4j/java_gateway.py in __call__(self, *args)
    1255         answer = self.gateway_client.send_command(command)
    1256         return_value = get_return_value(
-> 1257             answer, self.gateway_client, self.target_id, self.name)
    1258
    1259         for temp_arg in temp_args:

/opt/spark-2.4.3-bin-hadoop2.7/python/pyspark/sql/utils.py in deco(*a, **kw)
     61     def deco(*a, **kw):
     62         try:
---> 63             return f(*a, **kw)
     64         except py4j.protocol.Py4JJavaError as e:
     65             s = e.java_exception.toString()

/opt/spark-2.4.3-bin-hadoop2.7/python/lib/py4j-0.10.7-src.zip/py4j/protocol.py in get_return_value(answer, gateway_client, target_id, name)
    326         raise Py4JJavaError(
    327             "An error occurred while calling {0}{1}{2}.\n".
--> 328             format(target_id, ".", name), value)
    329     else:
    330         raise Py4JError(

Py4JJavaError: An error occurred while calling o146.json.
: java.lang.NoClassDefFoundError: org/apache/hadoop/fs/StorageStatistics
  at java.lang.Class.forName0(Native Method)
  at java.lang.Class.forName(Class.java:348)
  at org.apache.hadoop.conf.Configuration.getClassByNameOrNull(Configuration.java:2134)
  at org.apache.hadoop.conf.Configuration.getClassByName(Configuration.java:2099)
  at org.apache.hadoop.conf.Configuration.getClass(Configuration.java:2193)
  at org.apache.hadoop.fs.FileSystem.getFileSystemClass(FileSystem.java:2654)
  at org.apache.hadoop.fs.FileSystem.createFileSystem(FileSystem.java:2667)
  at org.apache.hadoop.fs.FileSystem.access$200(FileSystem.java:94)
  at org.apache.hadoop.fs.FileSystem$Cache.getInternal(FileSystem.java:2703)
  at org.apache.hadoop.fs.FileSystem$Cache.get(FileSystem.java:2685)

```

```

    at org.apache.hadoop.fs.FileSystem.get(FileSystem.java:373)
    at org.apache.hadoop.fs.Path.getFileSystem(Path.java:295)
    at org.apache.spark.sql.execution.streaming.FileStreamSink$.hasMeta
data(FileStreamSink.scala:45)
    at org.apache.spark.sql.execution.datasources.DataSource.resolveRel
ation(DataSource.scala:332)
    at org.apache.spark.sql.DataFrameReader.loadV1Source(DataFrameReade
r.scala:223)
    at org.apache.spark.sql.DataFrameReader.load(DataFrameReader.scala:
211)
    at org.apache.spark.sql.DataFrameReader.json(DataFrameReader.scala:
391)
    at sun.reflect.NativeMethodAccessorImpl.invoke0(Native Method)
    at sun.reflect.NativeMethodAccessorImpl.invoke(NativeMethodAccessor
Impl.java:62)
    at sun.reflect.DelegatingMethodAccessorImpl.invoke(DelegatingMethod
AccessorImpl.java:43)
    at java.lang.reflect.Method.invoke(Method.java:498)
    at py4j.reflection.MethodInvoker.invoke(MethodInvoker.java:244)
    at py4j.reflection.ReflectionEngine.invoke(ReflectionEngine.java:35
7)
    at py4j.Gateway.invoke(Gateway.java:282)
    at py4j.commands.AbstractCommand.invokeMethod(AbstractCommand.java:
132)
    at py4j.commands.CallCommand.execute(CallCommand.java:79)
    at py4j.GatewayConnection.run(GatewayConnection.java:238)
    at java.lang.Thread.run(Thread.java:748)
Caused by: java.lang.ClassNotFoundException: org.apache.hadoop.fs.StorageSt
atistics
    at java.net.URLClassLoader.findClass(URLClassLoader.java:382)
    at java.lang.ClassLoader.loadClass(ClassLoader.java:418)
    at java.lang.ClassLoader.loadClass(ClassLoader.java:351)
    ... 28 more

```

In [14]:

```

# read song data JSON file into data frame
df = spark.read.json(song_data, schema=songSchema)

```

In [15]:

```

# DEBUG
# songSchema given
df.printSchema()

```

```

root
|-- artist_id: string (nullable = true)
|-- artist_latitude: double (nullable = true)
|-- artist_location: string (nullable = true)
|-- artist_longitude: double (nullable = true)
|-- artist_name: string (nullable = true)
|-- duration: double (nullable = true)
|-- num_songs: integer (nullable = true)
|-- song_id: string (nullable = true)
|-- title: string (nullable = true)
|-- year: integer (nullable = true)

```


In [16]:

```
# df.count()
print('>> [' + str(df.count()) + '] songs from song_data read out in JSON-format')
```

```
>> [71] songs from song_data read out in JSON-format
```

In [17]:

```
# DEBUG
df.limit(5).toPandas()
```

Out[17]:

| | artist_id | artist_latitude | artist_location | artist_longitude | artist_name | dur |
|---|--------------------|-----------------|-------------------|------------------|---|-------|
| 0 | ARDR4AC1187FB371A1 | NaN | | NaN | Montserrat Caballé;Placido Domingo;Vicente Sar... | 511.1 |
| 1 | AREBBGV1187FB523D2 | NaN | Houston, TX | NaN | Mike Jones (Featuring CJ_Mello & Lil' Bran) | 173.6 |
| 2 | ARMAC4T1187FB3FA4C | 40.82624 | Morris Plains, NJ | -74.47995 | The Dillinger Escape Plan | 207.7 |
| 3 | ARPNLO1187FB3D52F | 40.71455 | New York, NY | -74.00712 | Tiny Tim | 43.3 |
| 4 | ARDNS031187B9924F0 | 32.67828 | Georgia | -83.22295 | Tim Wilson | 186.4 |

In [18]:

```
# DEBUG
df1 = df.filter(df.title == 'Young Boy Blues')
df1.limit(5).toPandas()
```

Out[18]:

| | artist_id | artist_latitude | artist_location | artist_longitude | artist_name | durati |
|---|--------------------|-----------------|-----------------|------------------|--------------|---------|
| 0 | ARGSJW91187B9B1D6B | 35.21962 | North Carolina | -80.01955 | JennyAnyKind | 218.775 |

In [52]:

```
# read song data JSON file into data frame
dfSdLoc = spark.read.json(song_data)
```

In [53]:

```
# DEBUG
# default Schema at read JSON
dfSdLoc.printSchema()
```

```
root
|-- artist_id: string (nullable = true)
|-- artist_latitude: double (nullable = true)
|-- artist_location: string (nullable = true)
|-- artist_longitude: double (nullable = true)
|-- artist_name: string (nullable = true)
|-- duration: double (nullable = true)
|-- num_songs: long (nullable = true)
|-- song_id: string (nullable = true)
|-- title: string (nullable = true)
|-- year: long (nullable = true)
```

In [54]:

```
# DEBUG
dfSdLoc.limit(5).toPandas()
```

Out[54]:

| | artist_id | artist_latitude | artist_location | artist_longitude | artist_name | duration |
|---|--------------------|-----------------|-----------------|------------------|-----------------------|----------|
| 0 | ARKFYS91187B98E58F | NaN | | NaN | Jeff And Sheri Easter | 267.702 |
| 1 | ARGSJW91187B9B1D6B | 35.21962 | North Carolina | -80.01955 | JennyAnyKind | 218.775 |
| 2 | ARD7TVE1187B99BFB1 | NaN | California - LA | NaN | Casual | 218.931 |

In [19]:

```
# extract columns to create songs table
song_columns = ["song_id", "title", "artist_id", "year", "duration"]
songs_table = df.select(song_columns) \
    .dropDuplicates()
```

In [20]:

```
# DEBUG
songs_table.printSchema()6
```

```
root
|-- song_id: string (nullable = true)
|-- title: string (nullable = true)
|-- artist_id: string (nullable = true)
|-- year: integer (nullable = true)
|-- duration: double (nullable = true)
```

In [21]:

```
# DEBUG
songs_table.limit(5).toPandas()
```

Out[21]:

| | song_id | | title | artist_id | year | duration |
|---|--------------------|---------------------------------|--------------------|--------------------|------|-----------|
| 0 | SOHUOAP12A8AE488E9 | | Floating | ARD842G1187B997376 | 1987 | 491.12771 |
| 1 | SOKEJEJ12A8C13E0D0 | The Urgency (LP Version) | ARC43071187B990240 | | 0 | 245.21098 |
| 2 | SONHOTT12A8C13493C | Something Girls | AR7G5I41187FB4CE6C | 1982 | | 233.40363 |
| 3 | SOHKNRJ12A6701D1F8 | Drop of Rain | AR10USD1187B99F3F1 | | 0 | 189.57016 |
| 4 | SOOLYAZ12A6701F4A6 | Laws Patrolling (Album Version) | AREBBGV1187FB523D2 | | 0 | 173.66159 |

In [22]:

```
# [-] VERIFICAR @ AWS...
# write songs table to parquet files partitioned by year and artist
songs_table.write.partitionBy('year', 'artist_id') \
    .parquet(output_data + 'songs_data/songs_table.parquet', 'overwrite')
```

In []:

```
# ??? Como me puedo alistar el contenido del directorio, para ver todo lo que ha echo?
# .. Que lo hizo correctamente! ;-)
```

In [23]:

```
# extract columns to create artists table
artists_column = ["artist_id", "artist_name as name", "artist_location as location", \
    "artist_latitude as latitude", "artist_longitude as longitude"]
artists_table = df.selectExpr(artists_column).dropDuplicates()
```

In [24]:

```
# DEBUG
artists_table.printSchema()

root
 |-- artist_id: string (nullable = true)
 |-- name: string (nullable = true)
 |-- location: string (nullable = true)
 |-- latitude: double (nullable = true)
 |-- longitude: double (nullable = true)
```

In [25]:

```
# DEBUG
artists_table.limit(5).toPandas()
```

Out[25]:

| | artist_id | name | location | latitude | longitude |
|---|--------------------|-----------------|-----------------|-----------|------------|
| 0 | ARPNLO1187FB3D52F | Tiny Tim | New York, NY | 40.71455 | -74.00712 |
| 1 | ARXR32B1187FB57099 | Gob | | NaN | NaN |
| 2 | AROGWRA122988FEE45 | Christos Dantis | | NaN | NaN |
| 3 | ARBGXIG122988F409D | Steel Rain | California - SF | 37.77916 | -122.42005 |
| 4 | AREVWGE1187B9B890A | Bitter End | Noci (BA) | -13.44200 | -41.99520 |

In [26]:

```
# write artists table to parquet files
artists_table.write.parquet(output_data + 'artists_data/artists_table.parquet', 'overwrite')
```

In []:

```
# ??? Como me puedo alistar el contenido del directorio, para ver todo lo que ha echo?
# .. Que lo hizo correctamente! ;-)
```

Step 4: Process Log Data

(read out, define schemas, write as parquet)

In [27]:

```
print('>> log data processed')

>> log data processed
```

In []:

```
# AWS:
#def process_log_data(spark, input_data, output_data):

# get filepath to log data file [log_data/2018/11/2018-11-17-events.json]
log_data = input_data + 'log_data/**/*.json'
```

In [45]:

```
# 1) experiment with a subset of the files,
# 2) and with local data (extract before from /data/log-data.zip)
log_data = input_data + 'log_data/2018/11/*.json'
```

In [68]:

```
# EINMALIG >>
# 3) for df.join(df.song == song_df.title) ALL required
# song_data_loc = input_data_local + 'song-data.zip'
import zipfile
with zipfile.ZipFile("data/log-data.zip") as zip_ref:
    zip_ref.extractall("data/log_data_unzipped/")
    # /data/log_data_unzipped/*.json
```

In [28]:

```
# 3) and with local data (extract before from /data/log-data.zip)
log_data = input_data + 'log_data_unzipped/*.json'
```

In [29]:

```
# read log data file into data frame
df = spark.read.json(log_data)
```

In [30]:

```
# number of lines
# df.count()
print('>> [' + str(df.count()) + '] logs entries read IN, of JSON logs_data')
```

```
>> [8056] logs entries read IN, of JSON logs_data
```

In [31]:

```
# DEBUG
df.printSchema()
```

```
root
|-- artist: string (nullable = true)
|-- auth: string (nullable = true)
|-- firstName: string (nullable = true)
|-- gender: string (nullable = true)
|-- itemInSession: long (nullable = true)
|-- lastName: string (nullable = true)
|-- length: double (nullable = true)
|-- level: string (nullable = true)
|-- location: string (nullable = true)
|-- method: string (nullable = true)
|-- page: string (nullable = true)
|-- registration: double (nullable = true)
|-- sessionId: long (nullable = true)
|-- song: string (nullable = true)
|-- status: long (nullable = true)
|-- ts: long (nullable = true)
|-- userAgent: string (nullable = true)
|-- userId: string (nullable = true)
```

In [32]:

```
# DEBUG
df.limit(5).toPandas()
```

Out[32]:

| | artist | auth | firstName | gender | itemInSession | lastName | length | level | location |
|---|-------------|-----------|-----------|--------|---------------|----------|-----------|-------|---------------------------------------|
| 0 | Harmonia | Logged In | Ryan | M | 0 | Smith | 655.77751 | free | San Jose-Sunnyvale-Santa Clara, CA |
| 1 | The Prodigy | Logged In | Ryan | M | 1 | Smith | 260.07465 | free | San Jose-Sunnyvale-Santa Clara, CA |
| 2 | Train | Logged In | Ryan | M | 2 | Smith | 205.45261 | free | San Jose-Sunnyvale-Santa Clara, CA |
| 3 | None | Logged In | Wyatt | M | 0 | Scott | NaN | free | Eureka-Arcata-Fortuna, CA |
| 4 | None | Logged In | Austin | M | 0 | Rosales | NaN | free | New York-Newark-Jersey City, NY-NJ-PA |

In [33]:

```
# filter by actions for song plays
df = df.filter(df.page == 'NextSong')
```

In [34]:

```
# number of reduce amount of lines
# df.count()
print('>> [' + str(df.count()) + '] songs filtered "NextSong" of logs_data')
```

```
>> [6820] songs filtered "NextSong" of logs_data
```

In [35]:

```
# extract columns for USERS table
# artists_table = # TYPO at resources (*.zip)!
users_columns = ["userId as user_id", "firstName as first_name", \
                 "lastName as last_name", "gender", "level"]
users_table = df.selectExpr(users_columns).dropDuplicates()
```

In [36]:

```
# DEBUG
users_table.printSchema()
```

```
root
|-- user_id: string (nullable = true)
|-- first_name: string (nullable = true)
|-- last_name: string (nullable = true)
|-- gender: string (nullable = true)
|-- level: string (nullable = true)
```

In [37]:

```
# DEBUG
users_table.limit(5).toPandas()
```

Out[37]:

| | user_id | first_name | last_name | gender | level |
|---|---------|------------|-----------|--------|-------|
| 0 | 26 | Ryan | Smith | M | free |
| 1 | 7 | Adelyn | Jordan | F | free |
| 2 | 71 | Ayleen | Wise | F | free |
| 3 | 81 | Sienna | Colon | F | free |
| 4 | 87 | Dustin | Lee | M | free |

In [38]:

```
# write USERS table to parquet files
users_table.write.parquet(output_data + 'users_data/users_table.parquet', 'overwrite')
```

In [40]:

```
time_columns = ["ts"]
dfTime = df.selectExpr(time_columns)
dfTime.printSchema()
```

```
root
|-- ts: long (nullable = true)
```

In [41]:

```
# DEBUG
dfTime.limit(2).toPandas()
```

Out[41]:

| | ts |
|---|---------------|
| 0 | 1542241826796 |
| 1 | 1542242481796 |

In [42]:

```
# SOLUTION:
get_datetime = udf(lambda x: str(datetime.fromtimestamp(int(x)/1000.0)) )
dfTimestamp = dfTime.withColumn("start_time", get_datetime(dfTime.ts))

dfTimestamp.printSchema()
dfTimestamp.show(3)
dfTimestamp.limit(3).toPandas()
```

```
root
|-- ts: long (nullable = true)
|-- start_time: string (nullable = true)
```

```
+-----+-----+
|          ts|          start_time|
+-----+-----+
|1542241826796|2018-11-15 00:30:...|
|1542242481796|2018-11-15 00:41:...|
|1542242741796|2018-11-15 00:45:...|
+-----+-----+
only showing top 3 rows
```

Out[42]:

| | ts | start_time |
|---|---------------|----------------------------|
| 0 | 1542241826796 | 2018-11-15 00:30:26.796000 |
| 1 | 1542242481796 | 2018-11-15 00:41:21.796000 |
| 2 | 1542242741796 | 2018-11-15 00:45:41.796000 |

In [43]:

```
# extract columns to create time table
time_table = dfTimestamp.select("start_time").dropDuplicates() \
    .withColumn("hour", hour(col("start_time"))) \
    .withColumn("day", dayofmonth(col("start_time"))) \
    .withColumn("week", weekofyear(col("start_time"))) \
    .withColumn("month", month(col("start_time"))) \
    .withColumn("year", year(col("start_time"))) \
    .withColumn("weekday", date_format(col("start_time"), 'E'))
```


In [44]:

```
# DEBUG
time_table.printSchema()
time_table.limit(5).toPandas()
```

```
root
|-- start_time: string (nullable = true)
|-- hour: integer (nullable = true)
|-- day: integer (nullable = true)
|-- week: integer (nullable = true)
|-- month: integer (nullable = true)
|-- year: integer (nullable = true)
|-- weekday: string (nullable = true)
```

Out[44]:

| | start_time | hour | day | week | month | year | weekday |
|---|----------------------------|------|-----|------|-------|------|---------|
| 0 | 2018-11-15 11:22:06.796000 | 11 | 15 | 46 | 11 | 2018 | Thu |
| 1 | 2018-11-15 18:09:32.796000 | 18 | 15 | 46 | 11 | 2018 | Thu |
| 2 | 2018-11-15 18:59:14.796000 | 18 | 15 | 46 | 11 | 2018 | Thu |
| 3 | 2018-11-15 19:01:55.796000 | 19 | 15 | 46 | 11 | 2018 | Thu |
| 4 | 2018-11-21 03:57:19.796000 | 3 | 21 | 47 | 11 | 2018 | Wed |

In [45]:

```
# write time table to parquet files partitioned by year and month
time_table.write.partitionBy("year", "month") \
    .parquet(output_data + 'time_data/time_table.parquet', 'overwrite')
```

Preparation and Code for songplays_table

In [46]:

```
get_datetime = udf(lambda x: str(datetime.fromtimestamp(int(x)/1000.0)))
df = df.withColumn("start_time", get_datetime(df.ts))
df = df.withColumn("year", year(col("start_time")))
df = df.withColumn("month", month(col("start_time")))
```

In [47]:

```
# DEBUG  
df.printSchema()  
df.limit(10).toPandas()
```


root

```
|-- artist: string (nullable = true)
|-- auth: string (nullable = true)
|-- firstName: string (nullable = true)
|-- gender: string (nullable = true)
|-- itemInSession: long (nullable = true)
|-- lastName: string (nullable = true)
|-- length: double (nullable = true)
|-- level: string (nullable = true)
|-- location: string (nullable = true)
|-- method: string (nullable = true)
|-- page: string (nullable = true)
|-- registration: double (nullable = true)
|-- sessionId: long (nullable = true)
|-- song: string (nullable = true)
|-- status: long (nullable = true)
|-- ts: long (nullable = true)
|-- userAgent: string (nullable = true)
|-- userId: string (nullable = true)
|-- start_time: string (nullable = true)
|-- year: integer (nullable = true)
|-- month: integer (nullable = true)
```

Out[47]:

| | artist | auth | firstName | gender | itemInSession | lastName | length | level | location |
|---|------------------------------------|-----------|-----------|--------|---------------|----------|-----------|-------|--------------------------------------|
| 0 | Harmonia | Logged In | Ryan | M | 0 | Smith | 655.77751 | free | San Jose-Sunnyvale-Santa Clara, CA |
| 1 | The Prodigy | Logged In | Ryan | M | 1 | Smith | 260.07465 | free | San Jose-Sunnyvale-Santa Clara, CA |
| 2 | Train | Logged In | Ryan | M | 2 | Smith | 205.45261 | free | San Jose-Sunnyvale-Santa Clara, CA |
| 3 | Sony Wonder | Logged In | Samuel | M | 0 | Gonzalez | 218.06975 | free | Houston-The Woodlands-Sugar Land, TX |
| 4 | Van Halen | Logged In | Tegan | F | 2 | Levine | 289.38404 | paid | Portland-South Portland, ME |
| 5 | Magic Sam | Logged In | Tegan | F | 3 | Levine | 132.04853 | paid | Portland-South Portland, ME |
| 6 | Edward Sharpe & The Magnetic Zeros | Logged In | Tegan | F | 4 | Levine | 306.31138 | paid | Portland-South Portland, ME |
| 7 | Usher featuring will.i.am | Logged In | Tegan | F | 5 | Levine | 395.72853 | paid | Portland-South Portland, ME |
| 8 | Helen Reddy | Logged In | Tegan | F | 7 | Levine | 176.50893 | paid | Portland-South Portland, ME |
| 9 | Taylor Swift | Logged In | Tegan | F | 8 | Levine | 201.06404 | paid | Portland-South Portland, ME |

10 rows × 21 columns



In [48]:

```
# df.count()
print('>> [' + str(df.count()) + '] songs filtered "NextSong" of logs_data')
```

```
>> [6820] songs filtered "NextSong" of logs_data
```

In [49]:

```
# read in song data to use for songplays table
song_df = spark.read.option("mergeSchema", "true").parquet(output_data + "songs_data/songs_table.parquet")
```

In [50]:

```
#song_df.count()
print('>> [' + str(song_df.count()) + '] songs readout from songs_table.PARQUET')
```

```
>> [71] songs readout from songs_table.PARQUET
```

In [51]:

```
# DEBUG
song_df.printSchema()
song_df.limit(5).toPandas()
```

```
root
 |-- song_id: string (nullable = true)
 |-- title: string (nullable = true)
 |-- duration: double (nullable = true)
 |-- year: integer (nullable = true)
 |-- artist_id: string (nullable = true)
```

Out[51]:

| | song_id | title | duration | year | artist_id |
|---|--------------------|---|-----------|------|--------------------|
| 0 | SOAOIBZ12AB01815BE | I Hold Your Hand In Mine [Live At Royal Albert... | 43.36281 | 2000 | ARPNLO1187FB3D52F |
| 1 | SONYPOM12A8C13B2D7 | I Think My Wife Is Running Around On Me (Taco ... | 186.48771 | 2005 | ARDNS031187B9924F0 |
| 2 | SODREIN12A58A7F2E5 | A Whiter Shade Of Pale (Live @ Fillmore West) | 326.00771 | 0 | ARLTWXK1187FB5A3F8 |
| 3 | SOYMRWW12A6D4FAB14 | The Moon And I (Ordinary Day Album Version) | 267.70240 | 0 | ARKFYS91187B98E58F |
| 4 | SOWQTQZ12A58A7B63E | Streets On Fire (Explicit Album Version) | 279.97995 | 0 | ARPFHN61187FB575F6 |

In [81]:

```
# DEBUG - JOIN - Testing (1/3)
df1 = song_df.filter(song_df.title == 'Setanta matins')
df1.limit(5).toPandas()
```

Out[81]:

| | song_id | title | duration | year | artist_id |
|---|--------------------|----------------|-----------|------|--------------------|
| 0 | SOZCTXZ12AB0182364 | Setanta matins | 269.58322 | 0 | AR5KOSW1187FB35FF4 |

In [82]:

```
# DEBUG - JOIN - Testing (2/3)
df2 = df.filter(df.song == 'Setanta matins') # Riverside / Young Boy Blues
df2.limit(5).toPandas()
```

Out[82]:

| | artist | auth | firstName | gender | itemInSession | lastName | length | level | location | me |
|---|--------|-----------|-----------|--------|---------------|----------|-----------|-------|------------------------------------|----|
| 0 | Elena | Logged In | Lily | F | 5 | Koch | 269.58322 | paid | Chicago-Naperville-Elgin, IL-IN-WI | |

In [92]:

```
# DEBUG - JOIN - Testing (3/3)
df3 = df.join(song_df, song_df.title == df.song ).select(df.song.alias("song (df)"), song_df.title.alias("title (song_df)"))
df3.count()
df3.limit(5).toPandas()
```

Out[92]:

| | song (df) | title (song_df) |
|---|----------------|-----------------|
| 0 | Setanta matins | Setanta matins |
| 1 | Intro | Intro |
| 2 | Intro | Intro |
| 3 | Intro | Intro |

In [52]:

```
# extract columns from joined song and log datasets to create songplays table
songplays_table = song_df.join(df, (song_df.title == df.song) ) \
    .select('start_time', \
            df.year, \
            df.month, \
            col('userId').alias("user_id"), \
            df.level, \
            song_df.song_id, \
            song_df.artist_id, \
            col('sessionId').alias("session_id"), \
            df.location, \
            col('userAgent').alias("user_agent") \
            )
```

In [53]:

```
songplays_table = songplays_table.withColumn("songplay_id", monotonically_increasing_id())
```

In [54]:

```
#songplays_table.count()
print('>> [' + str(songplays_table.count()) + '] songs found, on JOIN matching for songs
plays_table')
```

```
>> [4] songs found, on JOIN matching for songplays_table
```

In [55]:

```
# DEBUG
songplays_table.printSchema()
songplays_table.limit(3).toPandas()
```

```
root
|-- start_time: string (nullable = true)
|-- year: integer (nullable = true)
|-- month: integer (nullable = true)
|-- user_id: string (nullable = true)
|-- level: string (nullable = true)
|-- song_id: string (nullable = true)
|-- artist_id: string (nullable = true)
|-- session_id: long (nullable = true)
|-- location: string (nullable = true)
|-- user_agent: string (nullable = true)
|-- songplay_id: long (nullable = false)
```

Out[55]:

| | start_time | year | month | user_id | level | song_id | artist_id |
|---|-------------------------------|------|-------|---------|-------|--------------------|-------------------|
| 0 | 2018-11-21 21:56:47.796000 | 2018 | 11 | 15 | paid | SOZCTXZ12AB0182364 | AR5KOSW1187FB35FF |
| 1 | 2018-11-14 05:06:03.796000 | 2018 | 11 | 10 | free | SOGDBUF12A8C140FAA | AR558FS1187FB4565 |
| 2 | 2018-11-19 09:14:20.796000 | 2018 | 11 | 24 | paid | SOGDBUF12A8C140FAA | AR558FS1187FB4565 |

In [56]:

```
# write songplays table to parquet files partitioned by year and month
songplays_table.write.partitionBy("year", "month") \
    .parquet(output_data + 'songplays_data/songplays_table.parquet', 'overwri
te')
```

In [57]:

```
print('>> END! (main)')
```

```
>> END! (main)
```

Step 5: Run [etl.py] script

... with working clean code

In [58]:

```
!python 'etl.py' # to be executed at the Jupyter Notebook of the Udacity Workspace
```

```
>> Read Out Config-Infos from [dl.cfg]
>> START (main)
Ivy Default Cache set to: /root/.ivy2/cache
The jars for the packages stored in: /root/.ivy2/jars
:: loading settings :: url = jar:file:/opt/spark-2.4.3-bin-hadoop2.7/jars/ivy-2.4.0.jar!/org/apache/ivy/core/settings/ivysettings.xml
org.apache.hadoop#hadoop-aws added as a dependency
:: resolving dependencies :: org.apache.spark#spark-submit-parent-91ef8ed6-174b-40e9-ba64-9264353e7b01;1.0
    confs: [default]
    found org.apache.hadoop#hadoop-aws;2.10.0 in central
    found com.amazonaws#aws-java-sdk-bundle;1.11.271 in central
    found org.apache.commons#commons-lang3;3.4 in central
:: resolution report :: resolve 697ms :: artifacts dl 13ms
    :: modules in use:
    com.amazonaws#aws-java-sdk-bundle;1.11.271 from central in [default]
t]
    org.apache.commons#commons-lang3;3.4 from central in [default]
    org.apache.hadoop#hadoop-aws;2.10.0 from central in [default]
-----
--
|          |          modules          || artifacts
|
|      conf      | number| search|dwnlded|evicted|| number|dwnlde
d|
-----
--
|      default   |   3   |   0   |   0   |   0   ||   3   |   0
|
-----
--
:: retrieving :: org.apache.spark#spark-submit-parent-91ef8ed6-174b-40e9-ba64-9264353e7b01
    confs: [default]
    0 artifacts copied, 3 already retrieved (0kB/25ms)
23/03/21 00:02:20 WARN NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Using Spark's default log4j profile: org/apache/spark/log4j-defaults.properties
Setting default log level to "WARN".
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel).
23/03/21 00:02:22 WARN Utils: Service 'SparkUI' could not bind on port 4040. Attempting port 4041.
23/03/21 00:02:22 WARN Utils: Service 'SparkUI' could not bind on port 4041. Attempting port 4042.
>> spark session created
>> processing song data
>> [71] songs from song_data read out in JSON-format
>> song data processed
>> processing log data
>> [8056] logs entries read IN, of JSON logs_data
>> [6820] songs filtered page("NextSong") of logs_data
>> [6820] songs filtered "NextSong" of logs_data
>> [71] songs readout from songs_table.PARQUET
>> [4] songs found, on JOIN matching for songsplays_table
>> log data processed
>> END!
```

Step 6: Clean up Resources

(EMR, Notebook, IAM-Role)

- Export Jupyter Notebook
- Terminate EMR Cluster
- Delete IAM-Role

In []: